



Digital Article / AI and Machine Learning

Responsible AI Is Becoming a Growth Strategy

A playbook for turning good governance into a competitive edge.

by Michael Wade and Jochen Wirtz

Published on HBR.org / July 21, 2026 / Reprint [H099BQ](#)



Illustration by Maarten Huizing

There is a central paradox at the heart of the AI: The same capabilities that let firms personalize at scale, automate service delivery, and generate operational insight are the ones that surveil, manipulate, and discriminate. In 2025, [Apitor](#), a robot toy maker, was fined for harvesting children's location data and sending it to a third party in China. [Sirius XM Radio](#) is facing a lawsuit over an AI hiring tool that allegedly screened out Black applicants using racially proxied

variables such as zip code and alma mater before any human reviewed a resume. Australian regulators determined that two retailers owned by the conglomerate Wesfarmers deployed facial recognition on shoppers without adequate notice.

We expect such incidents to become common. As generative AI, agentic systems, and service robots move deeper into customer-facing operations, the gap between technological capability and organizational accountability is widening. Trust, once assumed and rarely measured, has become a scarce currency. Companies that fail to earn it will face regulatory action, reputational damage, and defection to competitors that chose a more trustworthy path.

In our eight years of researching Corporate Digital Responsibility (CDR)—the discipline of governing how firms design, deploy, and refine data and AI responsibly across their ecosystem—and observing deployments across many service industries, we have concluded that CDR is fast becoming a source of competitive advantage, and this article shows how to build that capability.

The Changing Trust Landscape

The stakes for CDR are rising because the underlying technology has become consequential in new ways. Three converging waves of AI have reshaped the landscape.

The first wave is generative AI. Large language and multimodal models now produce text, images, code, and decisions at scale. They can write the marketing copy that targets a firm's most vulnerable customers, generate fake reviews indistinguishable from authentic ones, and produce legal, medical, or financial content without the accountability structures that govern licensed professionals doing the same work. Because these models derive their own logic from vast datasets rather

than from explicit rules, their decisions are often inscrutable even to their creators, which raises hard questions of accountability and bias. A [2024 study](#) found that AI-generated websites contained deceptive design by default, the models having absorbed manipulation from the data on which they were trained.

The second wave is agentic AI, meaning systems that pursue goals autonomously and execute sequences of actions rather than simply responding to prompts. These systems can book flights, execute trades, manage cloud infrastructure, and resolve customer service requests without human review at each step. [McKinsey](#) found that 80 percent of organizations have already encountered problematic behavior from their AI agents in live deployments. In one controlled [Anthropic experiment](#), an agent facing shutdown mined an executive's personal emails, identified a compromising matter, and sent threats to prevent the shutdown from proceeding. The unsettling point is that the underlying capability is now widely deployed in enterprise environments.

The third wave is physical AI in the form of service robots that move, sense, and act in the world. They are already operating in restaurants, hotels, hospitals, and retail. When AI operates physically, [a failure can lead to serious consequences](#), such as a costly accident, downtime, or a diagnostic error.


Together, these waves have elevated CDR from a specialized IT concern to a board-level priority. The challenge is to build governance that can detect AI-related trust risks before they cause harm.

If CDR matters so much, why do companies struggle with it? The simple answer is that doing the right thing often comes at a cost. Firms are constantly navigating trade-offs between performance (doing well) and

customer well-being (doing good). We call this the CDR Calculus, and it exists to make that trade-off explicit at a moment when the calculation is becoming materially less favorable. At its core, the CDR Calculus is a 2×2 framework that maps any AI-related decision against two axes: its impact on business performance and its impact on customer well-being. Most frameworks of this kind merely describe. This one forces leadership teams to stop debating AI ethics in the abstract and start making concrete choices about specific systems.

The Corporate Digital Responsibility (CDR) Calculus

Every AI initiative creates trade-offs between business performance and stakeholder well-being. The CDR Calculus helps leaders distinguish between practices that build lasting trust and those that create hidden risks.



Source: Michael Wade and Jochen Wirtz



See more HBR charts in [Data & Visuals](#) on HBR.org.

Every significant AI deployment sits somewhere in this matrix, and the question is whether you make the calculation deliberately or discover your position after a regulator calls or a video goes viral.

Failures sit in the negative-business, negative-customer quadrant, where poorly governed deployments harm everyone. McDonald's three-year partnership with IBM to run AI voice ordering ended in 2024 after viral videos showed the system adding 260 McNuggets to a single order. Roughly 85 percent accuracy sounds adequate until you recognize that one interaction in seven failed visibly, in front of the queue and the internet. Agentic AI is creating an entirely new category of failure, illustrated when Amazon's coding agent Kiro caused a 13-hour AWS outage in December 2025 after autonomously deleting and recreating a production environment, with no human checkpoint and consequences that propagated at machine speed.

Temptations occupy the high-business, negative-customer quadrant, generating near-term revenue while eroding trust. Amazon's Prime subscription is the canonical case. Signing up took one click, while cancellation required a multi-step sequence the company internally called the "Iliad Flow," with each added hurdle tracked for its revenue impact. The FTC's 2025 settlement reached \$2.5 billion, including \$1.5 billion in refunds to 35 million consumers. Generative AI makes this dynamic more potent, since systems trained on behavioral data can identify and exploit individual psychological vulnerabilities at a scale no human sales team could match. Uber faced a parallel FTC action in 2025 over subscription enrollment and cancellation friction engineered in much the same way.

Sacrifices occupy the positive-customer, negative-business quadrant, capturing the real costs of responsible practice. Apple's 2021 App Tracking Transparency framework required explicit consent before cross-app tracking, at an estimated \$10 billion loss in annual advertising revenue to Meta and a genuine near-term cost to Apple itself. Within a month, 96 percent of U.S. iPhone users opted out, a revealed preference that validated the strategic bet. Leaders must be honest about costs

in this quadrant, because when the trade-off eventually surfaces, as it always does, someone needs to have been given the authority and budget to make it deliberately.

The Sweet Spot is the positive-positive quadrant that every organization should design toward, and it requires deliberate engineering. Wendy's FreshAI voice ordering, built with Google, improved average drive-through time by 22 seconds while holding accuracy above a threshold McDonald's never reached, and without the visible failures that ended McDonald's pilot. Patagonia likewise turned supply-chain transparency into a differentiator, growing to \$1.5 billion in revenue with a reported 73 percent loyalty rate. The same logic applies to AI: firms that can show their models are fair, their agents are bounded, and their robots are safe and reliable will earn the customer's trust, and the regulatory latitude that less disciplined competitors lose.

The matrix clarifies the business case. Good CDR does more than reduce regulatory exposure, fines, and operational disruption, because it builds brand equity and customer loyalty that competitors cannot easily copy. Those gains are real, but so are the costs, which include revenue foregone from cross-selling and dynamic pricing, reduced personalization, and sustained investment in governance infrastructure and culture. The aim is not to avoid every temptation or to accept every sacrifice, but to move deliberately toward the sweet spot given your specific business model, customer base, and competitive context.

The CDR Playbook

Knowing where a system sits in the CDR Calculus is only useful if you can act on it. The following three stages turn the calculus into organizational capability, and each one builds on the last.

Stage 1: Know what you have.

You cannot evaluate a tool if you do not know it is running. The most dangerous AI in any organization is the AI that leadership is unaware of, and most companies undercount badly. Beyond the flagship models announced by technology leaders, there are algorithmic pricing tools, automated hiring screeners, customer service bots, agentic workflows handling procurement, and third-party AI embedded in vendor platforms. The 2025 [MIT AI Agent Index](#) found that among the AI agents operating at frontier levels of autonomy, only a handful disclosed any safety evaluations, with developers revealing far more about capabilities than about risks. The same transparency gap runs through most enterprises.

Three concrete steps build the foundation. The first is an AI system registry, a living inventory of every deployed model, its decision scope, its training-data lineage, and its named human owner, maintained with the rigor that regulated banks apply to model risk inventories. The second is a risk tier for each system. The EU AI Act's categories of unacceptable, high, limited, and minimal risk offer a useful starting structure, because a recommendation engine is a materially different governance challenge from an agent approving insurance claims or a robot working alongside customers. The third is a short set of published non-negotiables, the actions your AI systems are never permitted to take regardless of the performance incentive. If a leadership team cannot name five such constraints within an hour, it does not yet have a functioning CDR strategy.

Together, these steps let you locate each system in the calculus, separating the genuine Sweet Spot candidates from the Temptations and Failures hiding in plain sight.

Diagnostic question: Can we name every AI system currently making decisions that affect our customers, and do we know who is accountable for each?

Stage 2: Build oversight into the machine.

The most common error here is treating AI governance as a review that happens after systems are built. Responsible AI bolted on late is systematically weaker, because bias, manipulation, and operational risk are most traceable at the design and training stages rather than after a public failure. This stage is how you move systems out of the Temptations and Failures quadrants before customers ever encounter them.

Microsoft's AI Red Team ran 67 operations in 2024, probing every flagship model before release for adversarial inputs and foreseeable misuse. Most organizations have no equivalent function, and establishing even a small team with a mandate to break systems before launch is likely the single most productive CDR investment available at this stage. For agentic AI, McKinsey's guidance is to design for trust first and speed second. Bounded autonomy, in which agents hold the minimum permissions needed for a task, irreversible decisions require human authorization, and every action is logged for later reconstruction, is the prerequisite for trusting agents rather than a brake on them. For physical AI, such as service robots, the stage adds something digitally native firms have rarely needed, namely physical incident protocols and the operational authority to deactivate a malfunctioning robot without escalating to an IT function.

Diagnostic question: If our most consequential AI system made a harmful decision tomorrow, could we reconstruct exactly what it did, why it did it, and who is accountable?

Stage 3: Turn CDR into competitive signal.

Organizations that reach this stage have stopped treating CDR as a cost center and started treating it as something they can demonstrate credibly to customers, regulators, ecosystem partners, and capital providers. This is where the sweet spot becomes external and defensible rather than a private claim.

The primary mechanism is structured public transparency, not the vague assurance that a company's AI is responsible, but the specific, auditable disclosure of what its systems do and how they are governed. Microsoft's annual [Responsible AI Transparency Report](#) names model families, documents red-teaming operations, and describes specific incidents and organizational responses, functioning as a governance artifact that happens to be public. A company that publishes specifics cannot later claim ignorance. The second mechanism is compensation alignment, since CDR objectives become durable only when they sit on the same leadership dashboard as revenue and retention and influence the same incentives. What is tied to compensation gets managed consistently. The third mechanism is participation in external standard-setting, where the most sophisticated organizations help shape emerging standards rather than wait for regulators to define responsible AI for them. Engagement with the NIST AI Risk Management Framework, the EU AI Act consultations, and commitments such as the [Frontier AI Safety Commitments](#) gives early participants influence over the rules their competitors will subsequently have to meet.

Diagnostic question: Could a well-informed regulator, customer, or institutional investor review our public CDR disclosures and conclude that we govern AI more responsibly than our competitors?

...

The numbers describe the situation plainly. According to the [Thales Digital Trust Index 2026](#), 93% of IT leaders already use, deploy, or plan AI initiatives, yet only 23% of consumers trust companies to handle AI and their data responsibly. That gap is as much a revenue problem as a communications one, and it is widening. Only [one in five companies](#) has a mature governance model for autonomous agents even as agentic deployment accelerates sharply, and nearly [two-thirds of organizations](#) name security and risk, rather than regulatory uncertainty or technical limits, as the primary barrier to scaling agentic AI. The bottleneck is trust, and specifically the absence of governance capable of earning it.

This is the moment CDR was designed for. The three waves of AI are not arriving in sequence but simultaneously, across organizations that mostly lack the governance to manage any one of them well. The CDR Calculus gives leaders a vocabulary for making trade-offs explicit rather than implicit, and the three-stage maturity path gives them a practical sequence for building governance that scales with deployment. Neither is sufficient on its own, and together they are now necessary.

Most organizations cannot yet answer the three diagnostic questions raised above, and that is precisely where the opportunity lies. The firms that close this gap first, treating governance not as a constraint on AI ambition but as the condition under which AI ambition becomes commercially sustainable, will inherit the customers, the talent, and the market position that less disciplined competitors are quietly forfeiting.

This article was originally published online on July 21, 2026.



Michael Wade is a professor of strategy and digital at IMD and director of the TONOMUS Global Center for Digital and AI Transformation. He is a coauthor of [GAIN: Demystifying GenAI for Office and Home](#) and [Twin Transformation](#), both from 2025.



Jochen Wirtz is a professor of marketing and the vice dean, MBA programs at the NUS Business School at the National University of Singapore.